

Household-Power-Consumption

Introduction

With the advent of new gadgets and a push towards greater electrification projects globally, power consumption is rising globally.

Thus, we can also expect that household or residential power consumption is so on the rise. With greater access to global power consumption data, forecasting power consumption is an emerging challenge.

An accurate forecast can help both the consumer as well as the supplier side. For the consumer, a power forecast helps in financial planning as making more green choices overall. For the supplier, an accurate forecast will definitely help in supply regulation. Thus, such models can help to optimize the overall supply chain of the household power industry.

Project Goal:

Track the power consumption of individual households in almost real time.



Used Language and Libraries:

```
└─ Python
    ├── Numpy
    ├── Pandas
    ├── Matplotlib
    ├── Seaborn
    ├── Sklearn
    ├── datetime
    └── missingno
```

```
In [1]: # ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
from datetime import datetime

from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
```

Dataset

Measurements of electric power consumption in one household with a one-minute sampling rate over a period of almost 4 years. Different electrical quantities and some sub-metering values are available.

Data Set Characteristics:

Associated Tasks:

Regression

Data Set Information:

This archive contains 185711 (sample of original dataset) measurements gathered between December 2006 and November 2010 (47 months).

Notes:

1.(global_active_power*1000/60 - sub_metering_1 - sub_metering_2 - sub_metering_3) represents the active energy consumed every minute (in watt hour) in the household by electrical equipment not measured in sub-meterings 1, 2 and 3.

2.The dataset contains some missing values in the measurements. All calendar timestamps are present in the dataset but for some timestamps, the measurement values are missing: a missing value is represented by the absence of value between two consecutive semi-colon attribute separators. For instance, the dataset shows missing values on April 28, 2007.

Attribute Information:

1.date: Date in format dd/mm/yyyy

2.time: time in format hh:mm:ss

3.global_active_power: household global minute-averaged active power (in kilowatt)

4.global_reactive_power: household global minute-averaged reactive power (in kilowatt)

5.voltage: minute-averaged voltage (in volt)

6.global_intensity: household global minute-averaged current intensity (in ampere)

7.sub_metering_1: energy sub-metering No. 1 (in watt-hour of active energy). It corresponds to the kitchen, containing mainly a dishwasher, an oven and a microwave (hot plates are not electric but gas powered).

8.sub_metering_2: energy sub-metering No. 2 (in watt-hour of active energy). It corresponds to the laundry room, containing a washing-machine, a tumble-drier, a refrigerator and a light.

9.sub_metering_3: energy sub-metering No. 3 (in watt-hour of active energy). It corresponds to an electric water-heater and an air-conditioner.

```
In [3]: df = pd.read_csv('data/Household Power Consumption.txt', sep=';')
df.head(5)
```

```
Out[3]:
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	16/12/2006	17:24:00	4.216	0.418	234.840	18.400	0.000	1.000	1.000
1	16/12/2006	17:25:00	5.360	0.436	233.630	23.000	0.000	1.000	1.000
2	16/12/2006	17:26:00	5.374	0.498	233.290	23.000	0.000	2.000	1.000
3	16/12/2006	17:27:00	5.388	0.502	233.740	23.000	0.000	1.000	1.000
4	16/12/2006	17:28:00	3.666	0.528	235.680	15.800	0.000	1.000	1.000

```
In [4]: df.shape
```

```
Out[4]: (185711, 9)
```

```
In [5]: df.describe()
```

```
Out[5]:
```

Sub_metering_3	
count	185702.000000
mean	6.553521
std	8.469009
min	0.000000
25%	0.000000
50%	0.000000
75%	17.000000
max	20.000000

```
In [6]: df['Date'].describe()
```

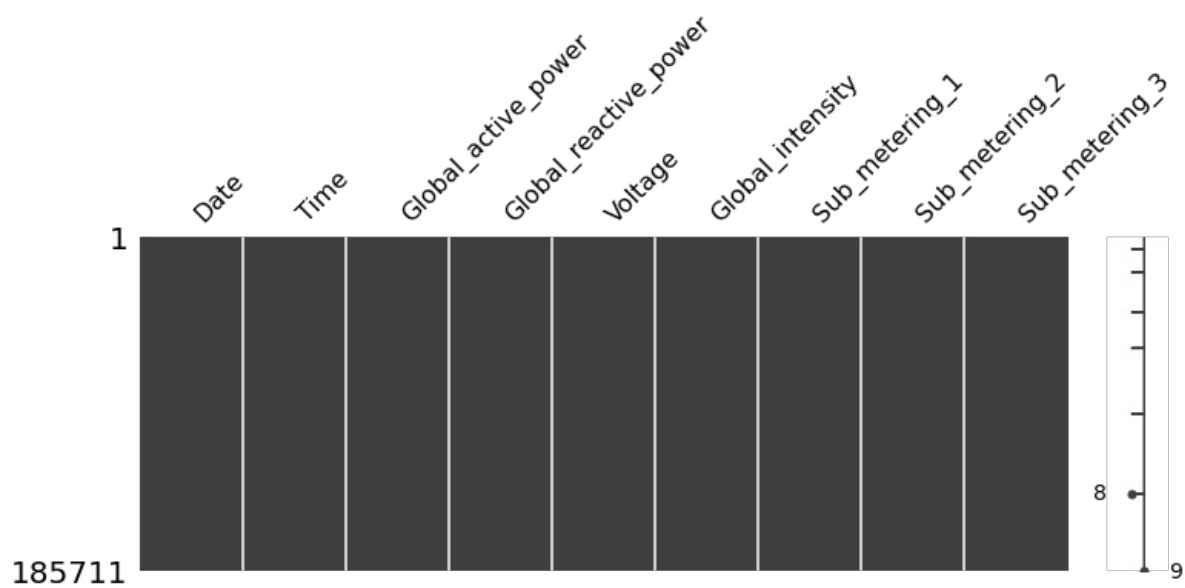
```
Out[6]: count      185711
unique        131
top      18/2/2007
freq         1440
Name: Date, dtype: object
```

```
In [7]: df.isna().sum()
```

```
Out[7]: Date      0
Time      0
Global_active_power  0
Global_reactive_power  0
Voltage      0
Global_intensity  0
Sub_metering_1  0
Sub_metering_2  0
Sub_metering_3  9
dtype: int64
```

```
In [8]: # Number of missing data.
msno.matrix(df, figsize=(12, 4))
```

```
Out[8]: <AxesSubplot:>
```



```
In [9]: # Replace all ? with NAN.
# This is useful in cases when you know the origin of the data and can be certain which values should be missing.
# Inplace=True parameter means that the changes will be made directly to the original DataFrame. (df), and it will not create a new DataFrame.
df.replace('?', np.nan, inplace=True)
df.replace(' ', np.nan, inplace=True)
df.replace(' ', np.nan, inplace=True)
df.replace(' ', np.nan, inplace=True)
```

```
In [10]: # 24*60 = 1440
df['Date'].value_counts()
```

```
Out[10]: 18/2/2007    1440
22/4/2007    1440
19/3/2007    1440
18/3/2007    1440
17/3/2007    1440
...
21/4/2007    1440
20/1/2007    1439
24/4/2007     995
16/12/2006    396
8             1
Name: Date, Length: 131, dtype: int64
```

```
In [11]: df.iloc[49850:49860,:]
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
49850	20/1/2007	08:14:00	1.122	0.106	237.880	4.800	0.000	0.000	0.000
49851	20/1/2007	08:15:00	1.122	0.108	238.400	4.600	0.000	0.000	0.000
49852	20/1/2007	08:16:00	1.118	0.108	238.230	4.600	0.000	0.000	0.000
49853	20/1/2007	08:17:00	1.606	0.102	237.830	7.600	0.000	0.000	0.000
49854	20/1/2007	08:18:00	3.124	0.084	236.420	13.200	0.000	0.000	0.000
49855	8	08:19:00	2.984	0.086	236.340	12.600	0.000	0.000	0.000
49856	20/1/2007	08:20:00	1.622	0.102	238.160	7.800	0.000	0.000	0.000
49857	20/1/2007	08:21:00	0.840	0.114	239.480	3.600	0.000	0.000	0.000
49858	20/1/2007	08:22:00	0.828	0.112	239.070	3.400	0.000	0.000	0.000
49859	20/1/2007	08:23:00	0.834	0.118	240.460	3.400	0.000	0.000	0.000

```
In [12]: df.replace({'8': '20/1/2007'}, inplace=True)
```

```
In [13]: df.iloc[49850:49860,:]
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
49850	20/1/2007	08:14:00	1.122	0.106	237.880	4.800	0.000	0.000	0.000
49851	20/1/2007	08:15:00	1.122	0.108	238.400	4.600	0.000	0.000	0.000
49852	20/1/2007	08:16:00	1.118	0.108	238.230	4.600	0.000	0.000	0.000
49853	20/1/2007	08:17:00	1.606	0.102	237.830	7.600	0.000	0.000	0.000
49854	20/1/2007	08:18:00	3.124	0.084	236.420	13.200	0.000	0.000	0.000
49855	20/1/2007	08:19:00	2.984	0.086	236.340	12.600	0.000	0.000	0.000
49856	20/1/2007	08:20:00	1.622	0.102	238.160	7.800	0.000	0.000	0.000
49857	20/1/2007	08:21:00	0.840	0.114	239.480	3.600	0.000	0.000	0.000
49858	20/1/2007	08:22:00	0.828	0.112	239.070	3.400	0.000	0.000	0.000
49859	20/1/2007	08:23:00	0.834	0.118	240.460	3.400	0.000	0.000	0.000

```
In [14]: df.iloc[49860:49870,:]
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
49860	20/1/2007	08:24:00	1.212	0.110	239.990	6.000	0.000	0.000	0.000
49861	20/1/2007	08:25:00	2.844	0.088	237.700	12.000	0.000	0.000	0.000
49862	20/1/2007	08:26:00	2.832	0.084	237.050	11.800	0.000	0.000	0.000
49863	20/1/2007	08:27:00	1.378	0.102	238.520	6.800	0.000	0.000	0.000
49864	20/1/2007	08:28:00	0.960	0.114	240.700	4.000	0.000	0.000	0.000
49865	20/1/2007	08:29:00	1.160	0.116	241.690	5.000	0.000	0.000	0.000
49866	20/1/2007	08:30:00	1.942	0.108	240.170	8.000	0.000	0.000	0.000
49867	20/1/2007	08:31:00	2.080	0.106	239.990	8.800	0.000	0.000	0.000
49868	20/1/2007	08:32:00	3.980	0.216	238.420	16.600	0.000	0.000	0.000
49869	20/1/2007	08:33:00	4.048	0.316	238.640	17.000	0.000	0.000	0.000

```
In [15]: # Check Duplicates
df["is_duplicate"] = df.duplicated()
df.shape
```

```
Out[15]: (185711, 10)
```

```

In [16]: # Remove Duplicates
df = df.drop_duplicates(subset=None,keep="first") # Dropping duplicate rows
df = df.drop(['is_duplicate'],axis=1) # Dropping created 'is_duplicate' column
df.shape

Out[16]: (185711, 9)

In [17]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 185711 entries, 0 to 185710
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date                                  185711 non-null object
1   Time                                  185711 non-null object
2   Global_active_power                  185701 non-null object
3   Global_reactive_power               185702 non-null object
4   Voltage                             185701 non-null object
5   Global_intensity                    185700 non-null object
6   Sub_metering_1                      185702 non-null object
7   Sub_metering_2                      185702 non-null object
8   Sub_metering_3                      185702 non-null float64
dtypes: float64(1), object(8)
memory usage: 14.2+ MB

In [18]: # make dataset numeric
df['Global_active_power'] = [float(x) for x in df['Global_active_power']]
df['Global_reactive_power'] = [float(x) for x in df['Global_reactive_power']]
df['Voltage'] = [float(x) for x in df['Voltage']]
df['Global_intensity'] = [float(x) for x in df['Global_intensity']]
df['Sub_metering_1'] = [float(x) for x in df['Sub_metering_1']]
df['Sub_metering_2'] = [float(x) for x in df['Sub_metering_2']]
df['Sub_metering_3'] = [float(x) for x in df['Sub_metering_3']]

In [19]: # Replace the missing values with the mean (avarage)
df['Global_active_power'].fillna(df['Global_active_power'].mean(),inplace=True)
df['Global_reactive_power'].fillna(df['Global_reactive_power'].mean(),inplace=True)
df['Voltage'].fillna(df['Voltage'].mean(),inplace=True)
df['Global_intensity'].fillna(df['Global_intensity'].mean(),inplace=True)
df['Sub_metering_1'].fillna(df['Sub_metering_1'].mean(),inplace=True)
df['Sub_metering_2'].fillna(df['Sub_metering_2'].mean(),inplace=True)
df['Sub_metering_3'].fillna(df['Sub_metering_3'].mean(),inplace=True)

In [20]: df['Date'].describe() # the unique values shows that there is a dataset of 130 day [from 2006-12-06 to 2007-04-
# and the top freq is 1440 [24 hours * 60 min] that means the data of each min a day

Out[20]: count      185711
unique        130
top          19/2/2007
freq         1440
Name: Date, dtype: object

In [21]: df.isna().sum()


Out[21]: Date      0
Time      0
Global_active_power  0
Global_reactive_power  0
Voltage      0
Global_intensity  0
Sub_metering_1      0
Sub_metering_2      0
Sub_metering_3      0
dtype: int64

In [22]: # Plot hourly Power Consumption for one day
df_plot = df[df.Date == '18/2/2007']
df_plot['hour'] = pd.to_datetime(df['Time'], format='%H:%M:%S').dt.hour

hourly_data = df_plot[['hour', 'Global_active_power']]
hourly_data = df_plot.groupby(['hour'], as_index = False)['Global_active_power'].agg('sum')

plt.plot(hourly_data['hour'], hourly_data['Global_active_power'], 'r', linewidth=2.0)
plt.xlabel('Hours')
plt.ylabel('Power Consumption')
plt.suptitle('One Day Power Consumption')
plt.grid(True)
plt.show()

```



```
In [23]: # merge Date and Time into a datetime
util_series = df['Date'] + ' ' + df['Time']
df['DateTime'] = util_series.apply(lambda x: datetime.strptime(x, '%d/%m/%Y %H:%M:%S'))
df.head(5)
```

```
Out[23]:
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	16/12/2006	17:24:00	4.216	0.418	234.84	18.4	0.0	1.0	1
1	16/12/2006	17:25:00	5.360	0.436	233.63	23.0	0.0	1.0	1
2	16/12/2006	17:26:00	5.374	0.498	233.29	23.0	0.0	2.0	1
3	16/12/2006	17:27:00	5.388	0.502	233.74	23.0	0.0	1.0	1
4	16/12/2006	17:28:00	3.666	0.528	235.68	15.8	0.0	1.0	1

```
In [24]: dt = df.copy()
```

```
In [25]: #drop old Date and Time columns
df.drop(columns = ['Date','Time'],inplace = True)
df
```

Out[25]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3	DateTime
0	4.216	0.418	234.84	18.4	0.0	1.0	17.0	2006-12-16 17:24:00
1	5.360	0.436	233.63	23.0	0.0	1.0	16.0	2006-12-16 17:25:00
2	5.374	0.498	233.29	23.0	0.0	2.0	17.0	2006-12-16 17:26:00
3	5.388	0.502	233.74	23.0	0.0	1.0	17.0	2006-12-16 17:27:00
4	3.666	0.528	235.68	15.8	0.0	1.0	17.0	2006-12-16 17:28:00
...
185706	0.320	0.082	235.19	1.6	0.0	2.0	0.0	2007-04-24 16:30:00
185707	0.328	0.078	234.53	1.6	0.0	1.0	0.0	2007-04-24 16:31:00
185708	0.344	0.088	234.90	1.8	0.0	1.0	0.0	2007-04-24 16:32:00
185709	0.416	0.174	235.29	2.0	0.0	1.0	0.0	2007-04-24 16:33:00
185710	0.408	0.174	234.82	2.0	0.0	1.0	0.0	2007-04-24 16:34:00

185711 rows × 8 columns

In [26]:

```
# Rearranging columns
cols = list(df.columns)
cols = cols[-1:] + cols[:-1]
data = df[cols]
data
```

Out[26]:


	DateTime	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	2006-12-16 17:24:00	4.216	0.418	234.84	18.4	0.0	1.0	17.0
1	2006-12-16 17:25:00	5.360	0.436	233.63	23.0	0.0	1.0	16.0
2	2006-12-16 17:26:00	5.374	0.498	233.29	23.0	0.0	2.0	17.0
3	2006-12-16 17:27:00	5.388	0.502	233.74	23.0	0.0	1.0	17.0
4	2006-12-16 17:28:00	3.666	0.528	235.68	15.8	0.0	1.0	17.0
...
185706	2007-04-24 16:30:00	0.320	0.082	235.19	1.6	0.0	2.0	0.0
185707	2007-04-24 16:31:00	0.328	0.078	234.53	1.6	0.0	1.0	0.0
185708	2007-04-24 16:32:00	0.344	0.088	234.90	1.8	0.0	1.0	0.0
185709	2007-04-24 16:33:00	0.416	0.174	235.29	2.0	0.0	1.0	0.0
185710	2007-04-24 16:34:00	0.408	0.174	234.82	2.0	0.0	1.0	0.0

185711 rows × 8 columns

Data Analysis

```
In [27]: plt.figure(figsize=(40,10))  
  
# plot for 5 days = 7200 minutes  
plt.plot(data['DateTime'][:7200],[float(x) for x in data['Global_active_power']][:7200])
```


```
Out[27]: [<matplotlib.lines.Line2D at 0x2f5d2e562b0>]
```




```
In [28]: data.isnull().any(axis = 1).sum()
```

```
Out[28]: 0
```

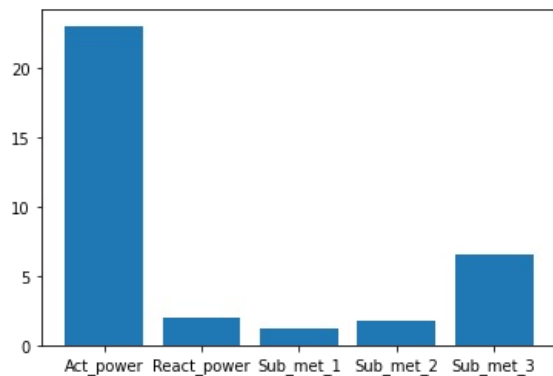
```
In [29]: # The correlations between 'Global_intensity', 'Global_active_power'  
sns.jointplot(x='Global_intensity', y='Global_active_power', data=df)  
plt.show()
```



```
In [30]: # The correlations between 'Voltage' and 'Global_active_power'  
  
sns.jointplot(x='Voltage', y='Global_active_power', data=df)  
plt.show()
```





```
In [31]: # rename columns
names = ['Act_power', 'React_power', 'Sub_met_1', 'Sub_met_2', 'Sub_met_3']
values = [23.042456, 2.027712, 1.232561, 1.783675, 6.553521]
plt.bar(names, values)
plt.show()
```



From above two plots it is seen that 'Global_intensity' and 'Global_active_power' correlated. But 'Voltage', 'Global_active_power' are less correlated. This is important observation for machine learning purpose.

```
In [32]: # Calculate correlations among Columns
#There is a strong correlation between 'Global_active_power' and 'Global_intensity' beacuse  $P = V * I$ 
plt.figure(figsize=(10,10))
ax=sns.heatmap(df.corr(), vmin=-1, cmap='coolwarm', annot=True);
ax.set_ylim(8, 0)
plt.show()
```



```
In [33]: def month(x):
          return x.month
data['Month']= data['DateTime'].apply(lambda x: month(x))

def Day(x):
    return x.Day
data['Day']=data['DateTime'].apply(lambda x: x.day)

def hour(x):
    return x.hour
data['Hour']= data['DateTime'].apply(lambda x: hour(x))
```


```
In [34]: data.drop('DateTime',axis=1,inplace=True)
data.tail()
```

```
Out[34]:
```

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3	Month	Day
185706	0.320	0.082	235.19	1.6	0.0	2.0	0.0	4	24
185707	0.328	0.078	234.53	1.6	0.0	1.0	0.0	4	24
185708	0.344	0.088	234.90	1.8	0.0	1.0	0.0	4	24
185709	0.416	0.174	235.29	2.0	0.0	1.0	0.0	4	24
185710	0.408	0.174	234.82	2.0	0.0	1.0	0.0	4	24


```
In [35]: # Monthly change of Global_active_power
plt.figure(figsize=(5,5))
plt.grid('both')
sns.lineplot(x="Month", y="Global_active_power",data=data)
```

```
Out[35]: <AxesSubplot:xlabel='Month', ylabel='Global_active_power'>
```




```
In [36]: # Daily change of Global_active_power
plt.figure(figsize=(5,5))
plt.grid('both')
sns.lineplot(x="Day", y="Global_active_power", data=data)
```

```
Out[36]: <AxesSubplot:xlabel='Day', ylabel='Global_active_power'>
```




```
In [37]: plt.figure(figsize=(5,5))
plt.grid('x')
sns.barplot(x="Month", y="Global_active_power", data=data)
```

```
Out[37]: <AxesSubplot:xlabel='Month', ylabel='Global_active_power'>
```



```
In [38]: plt.figure(figsize=(8,8))
sns.barplot(x="Day", y="Global_active_power", data=data)
```

```
Out[38]: <AxesSubplot:xlabel='Day', ylabel='Global_active_power'>
```



Data Pre-processing

```
In [39]: dt.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	16/12/2006	17:24:00	4.216	0.418	234.84	18.4	0.0	1.0	17.0
1	16/12/2006	17:25:00	5.360	0.436	233.63	23.0	0.0	1.0	16.0
2	16/12/2006	17:26:00	5.374	0.498	233.29	23.0	0.0	2.0	17.0
3	16/12/2006	17:27:00	5.388	0.502	233.74	23.0	0.0	1.0	17.0
4	16/12/2006	17:28:00	3.666	0.528	235.68	15.8	0.0	1.0	17.0

```
In [40]: X = df[['Global_reactive_power','Voltage','Global_intensity','Sub_metering_1','Sub_metering_2','Sub_metering_3']
y = df['Global_active_power']
```

```
In [41]: X.head()
```

	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	0.418	234.84	18.4	0.0	1.0	17.0
1	0.436	233.63	23.0	0.0	1.0	16.0
2	0.498	233.29	23.0	0.0	2.0	17.0
3	0.502	233.74	23.0	0.0	1.0	17.0
4	0.528	235.68	15.8	0.0	1.0	17.0

```
In [42]: dt.shape
```

```
Out[42]: (185711, 10)
```

```
In [43]: # Rearranging columns
cols = list(dt.columns)
cols = cols[-1:] + cols[:-1]
dt2 = dt[cols]
dt2.head(5)
```

Out[43]:	DateTime	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	S
0	2006-12-16 17:24:00	16/12/2006	17:24:00	4.216	0.418	234.84	18.4	0.0	1.0	
1	2006-12-16 17:25:00	16/12/2006	17:25:00	5.360	0.436	233.63	23.0	0.0	1.0	
2	2006-12-16 17:26:00	16/12/2006	17:26:00	5.374	0.498	233.29	23.0	0.0	2.0	
3	2006-12-16 17:27:00	16/12/2006	17:27:00	5.388	0.502	233.74	23.0	0.0	1.0	
4	2006-12-16 17:28:00	16/12/2006	17:28:00	3.666	0.528	235.68	15.8	0.0	1.0	

```
In [44]: # Column for for the remainder of sub metering
# Calculate the energy consumption in kilowatt-hours
eq1 = (dt2['Global_active_power'] * 1000/60) # Convert Global_active_power to kilowatt-hours

# Calculate the sum of sub metering 1, 2, and 3
eq2 = dt2['Sub_metering_1'] + dt2['Sub_metering_2'] + dt2['Sub_metering_3']

# Calculate the value for the new column 'Sub_metering_4'
dt2['Sub_metering_4'] = eq1 - eq2

# Display the first 3 rows of the updated DataFrame
dt2.head(3)
```

Out[44]:	DateTime	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	S
0	2006-12-16 17:24:00	16/12/2006	17:24:00	4.216	0.418	234.84	18.4	0.0	1.0	
1	2006-12-16 17:25:00	16/12/2006	17:25:00	5.360	0.436	233.63	23.0	0.0	1.0	
2	2006-12-16 17:26:00	16/12/2006	17:26:00	5.374	0.498	233.29	23.0	0.0	2.0	

```
In [45]: # drop old Date and Time columns
dt2.drop(columns = ['Date','Time'],inplace = True)
dt2.head(3)
```

Out[45]:	DateTime	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3	Sub_r
0	2006-12-16 17:24:00	4.216	0.418	234.84	18.4	0.0	1.0		17.0
1	2006-12-16 17:25:00	5.360	0.436	233.63	23.0	0.0	1.0		16.0
2	2006-12-16 17:26:00	5.374	0.498	233.29	23.0	0.0	2.0		17.0

Splitting dataset into train, validation, and test

```
In [46]: # First, split the data into testing 30% and the rest (training + validation) 70%
X_rem, X_test, y_rem, y_test = train_test_split(
    X,
    y,
    random_state=0,
    test_size=0.30,
)

# Second, split the remaining data (training + validation) into validation 30%
X_train, X_val, y_train, y_val = train_test_split(
    X_rem,
    y_rem,
    random_state=0,
    test_size=0.30,
)

print("x_train", X_train.shape)
print("y_train", y_train.shape)
```

```
print("-----")
print("x_test", X_test.shape)
print("y_test", y_test.shape)
print("-----")
print("x_val", X_val.shape)
print("y_val", y_val.shape)
```

```
x_train (90997, 6)
y_train (90997,)
-----
x_test (55714, 6)
y_test (55714,)
-----
x_val (39000, 6)
y_val (39000,)
```

Helper Function : Evaluate Model

```
In [47]: def eval_model(model, X_train, y_train, _val, y_val):
        model.fit(X_train, y_train)
        y_pred_train = model.predict(X_train)
        y_pred_val = model.predict(X_val)
        r2_train = r2_score(y_train, y_pred_train)
        r2_val = r2_score(y_val, y_pred_val)
        return r2_train, r2_val
```

Feature scaling


```
In [48]: scaler = StandardScaler()
        scaler.fit(X_train)

        X_train_scaled = scaler.transform(X_train)
        X_val_scaled = scaler.transform(X_val)
        X_test_scaled = scaler.transform(X_test)
```

Linear Regression

```
In [49]: lm = LinearRegression()
        eval_model(lm, X_train_scaled, y_train, X_val_scaled, y_val)
        predictions = lm.predict(X_test_scaled)
```

```
In [50]: plt.figure(figsize=(5,5))
        plt.scatter(y_test, predictions)
        plt.xlabel('Y Test')
        plt.ylabel('Predicted Y')
        plt.show()
```




```
In [51]: print('R2 for Linear Regression:', metrics.r2_score(y_test, predictions))
```

```
R2 for Linear Regression: 0.9985845710955411
```

```
In [52]: # create a distribution plot (histogram) of the differences between the actual values (y_test) and the predicted values
        # The resulting array represents the residuals, which are the errors made by the model in its predictions
        #
        sns.distplot((y_test-predictions),bins=50)
```

```
Out[52]: <AxesSubplot:xlabel='Global_active_power', ylabel='Density'>
```



```
In [53]: df_result = pd.DataFrame({'Actual Value': y_test, 'Predicted Value': predictions})
df_result
```

```
Out[53]:
```

	Actual Value	Predicted Value
15507	1.270	1.265892
96436	1.804	1.765216
139278	0.302	0.260394
151520	1.982	1.975361
38959	2.376	2.328261
...
104178	0.216	0.187842
13129	1.396	1.363891
77756	2.008	1.971718
88938	1.314	1.309561
95409	1.424	1.378481

55714 rows × 2 columns

Non Linear Regression

Hyper Parameters Tuning for degree in Polynomial Features

```
In [54]: degree_values = list(range(1, 5)) # 1 2 3 4
r2_train_values = []
r2_val_values = []

for degree in degree_values:
    nonlinear_reg = Pipeline([
        ('feature_scaling', StandardScaler()),
        ('polynomial_features', PolynomialFeatures(degree=degree)),
        ('regression', LinearRegression()),
    ])
    r2_train, r2_val = eval_model(nonlinear_reg, X_train_scaled, y_train, X_val_scaled, y_val)
    r2_train_values.append(r2_train)
    r2_val_values.append(r2_val)

results = pd.DataFrame({
    'degree': degree_values,
    'r2_train': r2_train_values,
    'r2_val': r2_val_values,
})

results
```

```
Out[54]:
```

	degree	r2_train	r2_val
0	1	0.998498	-9.012111e+01
1	2	0.998795	-3.189084e+03
2	3	0.998898	-1.942921e+05
3	4	0.998953	-8.506383e+11

from above DataFrame we take R2 Score with degree 2

from above Data-frame we take R2-Scores with degree 3

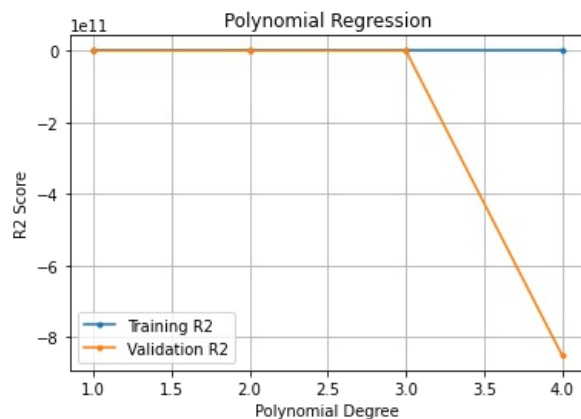
ploting degree in x-axis and r2-score (r2_train, r2_val) in y-axis

```
In [55]: import matplotlib.pyplot as plt

# Plotting the training R2 scores
plt.plot(results['degree'], results['r2_train'], label='Training R2', marker='.')

# Plotting the validation R2 scores
plt.plot(results['degree'], results['r2_val'], label='Validation R2', marker='.')

# Adding labels and title
plt.xlabel('Polynomial Degree')
plt.ylabel('R2 Score')
plt.title('Polynomial Regression')
plt.legend()
plt.grid('both')
# Show the plot
plt.show()
```




```
In [56]: best_model = Pipeline([
    ('feature_scaling', StandardScaler()),
    ('polynomial_features', PolynomialFeatures(degree=3)),
    ('regression', LinearRegression()),
])
best_model.fit(X_train_scaled, y_train)
y_pred_test = best_model.predict(X_test_scaled)
r2_score(y_test, y_pred_test)
```

Out[56]: 0.9989754556698835

Visualizing Our test predictions

```
In [57]: plt.figure(figsize=(5, 5))
plt.scatter(y_test, y_pred_test, alpha=0.4)
plt.plot([0, 10], [0, 10], c='red')
plt.xlabel('y_true')
plt.ylabel('y_pred')
plt.show()
```



Random Forest (RF)


```
In [58]: from sklearn.ensemble import RandomForestRegressor
rf=RandomForestRegressor(n_estimators=100,random_state=0,max_depth=3)
rf.fit(X_train, y_train)
y_pred_train = rf.predict(X_train)
y_pred_val = rf.predict(X_val)
r2_train = r2_score(y_train, y_pred_train)
r2_val = r2_score(y_val, y_pred_val)

print(r2_train, r2_val)
```

0.9796902922896603 0.9794722746614143

Optimizing n_estimators

```
In [59]: n_estimators_values = [5,10,15,20,25,30,35,40]

r2_train_values = []
r2_val_values = []

for n_estimators in n_estimators_values:
    model = RandomForestRegressor(n_estimators=n_estimators, max_depth=3, random_state=0)
    r2_train, r2_val = eval_model(model, X_train, y_train, X_val, y_val)
    r2_train_values.append(r2_train)
    r2_val_values.append(r2_val)


results = pd.DataFrame({
    'n_estimators': n_estimators_values,
    'r2_train': r2_train_values,
    'r2_val': r2_val_values,
})
```

```
In [60]: results['diff']=results['r2_train']-results['r2_val']
results
```

```
Out[60]:
```

	n_estimators	r2_train	r2_val	diff
0	5	0.979493	0.979168	0.000325
1	10	0.979567	0.979235	0.000333
2	15	0.979856	0.979553	0.000303
3	20	0.979935	0.979640	0.000295
4	25	0.980064	0.979767	0.000297
5	30	0.979749	0.979473	0.000275
6	35	0.979809	0.979520	0.000289
7	40	0.979843	0.979579	0.000263

```
In [61]: results.plot(x='n_estimators', y=['r2_train', 'r2_val'], figsize=(5,5))
plt.grid(axis='both')
plt.xticks(n_estimators_values)
plt.show()
```



Optimizing max_depth

```
In [62]: max_depth_values = [1, 2, 3, 4, 5, 6, 7, 8, 9]
r2_train_values = []
r2_val_values = []

for max_depth in max_depth_values:
    model = RandomForestRegressor(n_estimators=25, max_depth=max_depth, random_state=0)
    r2_train, r2_val = eval_model(model, X_train, y_train, X_val, y_val)
```

```

r2_train_values.append(r2_train)
r2_val_values.append(r2_val)

results = pd.DataFrame({
    'max_depth': max_depth_values,
    'r2_train': r2_train_values,
    'r2_val': r2_val_values,
})

results

```

Out[62]:

	max_depth	r2_train	r2_val
0	1	0.686862	0.686397
1	2	0.929951	0.928336
2	3	0.980064	0.979767
3	4	0.994503	0.994284
4	5	0.997886	0.997754
5	6	0.998523	0.998389
6	7	0.998763	0.998587
7	8	0.998970	0.998764
8	9	0.999113	0.998864

In [63]:

```

results['diff']=results['r2_train']-results['r2_val']
results

```

Out[63]:


	max_depth	r2_train	r2_val	diff
0	1	0.686862	0.686397	0.000466
1	2	0.929951	0.928336	0.001615
2	3	0.980064	0.979767	0.000297
3	4	0.994503	0.994284	0.000219
4	5	0.997886	0.997754	0.000132
5	6	0.998523	0.998389	0.000134
6	7	0.998763	0.998587	0.000177
7	8	0.998970	0.998764	0.000207
8	9	0.999113	0.998864	0.000249

In [64]:

```

results.plot(x='max_depth', y=['r2_train', 'r2_val'], figsize=(8, 6))
plt.grid(axis='both') # estimators for max_depth = 3
plt.xticks(max_depth_values)
plt.show()

```



In [65]:

```

model_best = RandomForestRegressor(n_estimators=100, max_depth=3, random_state=0)
model_best.fit(X_train, y_train)
y_pred_test = model_best.predict(X_test)
print(r2_score(y_test, y_pred_test))

```

0.9803752238470986

In [66]:

```

print(f'Training R2 Score: {r2_train}')
print(f'Validation R2 Score: {r2_val}')
print(f'Test R2 Score: {r2_score(y_test, y_pred_test)}')

```


Training R2 Score: 0.9991128401476651
Validation R2 Score: 0.9988640587661459
Test R2 Score: 0.9803752238470986

KNN


```
In [67]: from sklearn.neighbors import KNeighborsRegressor  
from sklearn.metrics import mean_squared_error
```

```
In [68]: # Create a KNN regressor  
model5 = KNeighborsRegressor(n_neighbors=5, weights='distance')  
  
# Evaluate the KNN regressor using the eval_model function  
r2_train, r2_val = eval_model(model5, X_train_scaled, y_train, X_val_scaled, y_val)
```


```
In [69]: # Plotting for Training Set  
y_train_pred = model5.predict(X_train_scaled)  
plt.scatter(y_train, y_train_pred, color='green', label='Actual vs Predicted (Training Set)')  
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)], linestyle='--', color='orange', label='Perf  
plt.xlabel('True Values (Training Set)')  
plt.ylabel('Predicted Values (Training Set)')  
plt.title('Actual vs Predicted Values for KNN Regression (Training Set)')  
plt.legend()  
plt.show()
```



```
In [70]: # Plotting for Validation Set  
y_val_pred = model5.predict(X_val_scaled)  
plt.scatter(y_val, y_val_pred, color='purple', label='Actual vs Predicted (Validation Set)')  
plt.plot([min(y_val), max(y_val)], [min(y_val), max(y_val)], linestyle='--', color='pink', label='Perfect Predi  
plt.xlabel('True Values (Validation Set)')  
plt.ylabel('Predicted Values (Validation Set)')  
plt.title('Actual vs Predicted Values for KNN Regression (Validation Set)')  
plt.legend()  
plt.show()
```



```
In [71]: # Plotting for Test Set  
y_test_pred = model5.predict(X_test_scaled)  
plt.scatter(y_test, y_test_pred, color='blue', label='Actual vs Predicted (Test Set)')  
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], linestyle='--', color='red', label='Perfect Pre  
plt.xlabel('True Values (Test Set)')  
plt.ylabel('Predicted Values (Test Set)')  
plt.title('Actual vs Predicted Values for KNN Regression (Test Set)')  
plt.legend()  
plt.show()
```



```
In [72]: print(f'Training R2 Score: {r2_train}')
print(f'Validation R2 Score: {r2_val}')
print(f'Test R2 Score: {r2_score(y_test, y_test_pred)}')
```

Training R2 Score: 0.999982555075161
 Validation R2 Score: -0.46667192273359026
 Test R2 Score: 0.9967624762663999

Decision Tree Regressor

```
In [73]: #DecisionTreeRegressor( max_depth=None, min_samples_split=2)
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV
dt_reg = DecisionTreeRegressor(random_state = 0, max_depth = 3)
dt_reg2 = DecisionTreeRegressor(random_state = 0, max_depth = 5)
dt_reg.fit(X_train, y_train)
dt_reg2.fit(X_train, y_train)
```

Out[73]: DecisionTreeRegressor(max_depth=5, random_state=0)

```
In [74]: print("train score : ", dt_reg.score(X_train, y_train))
```

train score : 0.9767084968301715

```
In [75]: print("test score : ", dt_reg.score(X_test, y_test))
```

test score : 0.977321881056382

```
In [76]: print("train score : ", dt_reg2.score(X_train, y_train))
```

train score : 0.9968573699979861

```
In [77]: print("test score : ", dt_reg2.score(X_test, y_test))
```

test score : 0.9969731151944378

```
In [78]: #caculate new predict
y_pred = dt_reg.predict(X_test)
```

```
In [79]: print("First 10 test values : \n",y_test[:5])
```

First 10 test values :
 15507 1.270
 96436 1.804
 139278 0.302
 151520 1.982
 38959 2.376
 Name: Global_active_power, dtype: float64

```
In [80]: print("First 10 predicted values : \n",y_pred[:5])
```

First 10 predicted values :
 [1.34432154 1.74082728 0.29937132 1.74082728 2.3861735]

```
In [81]: y_pred2 = dt_reg2.predict(X_test)
```

```
In [82]: print("First 10 predicted values : \n",y_pred2[:5])
```

First 10 predicted values :
 [1.31487324 1.78303658 0.28378695 1.89597413 2.42926898]

```
In [83]: # # Define the parameter grid
# param_grid_dt = {
#     'max_depth': [3, 5, 7, 9],
#     'min_samples_split': [2, 5, 10]
# }

# dt_reg = DecisionTreeRegressor()
```

```
# dt_reg = GridSearchCV(dt_reg, param_grid=param_grid_dt, cv=5, scoring='r2')
# dt_reg.fit(X_train, y_train)

# best_dt = dt_reg.best_estimator_
# best_score_dt = dt_reg.best_score_

# print("Best Decision Tree parameters:", best_dt)
# print("Best Decision Tree R2 score:", best_score_dt)
```

Support vector machine

```
In [84]: from sklearn.svm import SVR
```

```
In [85]: SVRModel = SVR(C=1.0, epsilon=0.1, kernel = 'rbf')
```

```
In [86]: train_r2, val_r2 = eval_model(SVRModel, X_train, y_train, X_val, y_val)
```

```
In [87]: print('R2 Score train equals ', train_r2)
print('R2 Score validation equals ', val_r2)
```

```
R2 Score train equals  0.9977184241687787
R2 Score validation equals  0.9976645499790809
```

```
In [88]: y_pred_test = SVRModel.predict(X_test)
```

```
In [89]: print(y_pred_test[:10])
```

```
[1.20122323  1.7326498  0.23313083  1.91876723  2.29039761  0.48460688
 0.30645171  2.17763018  2.51502357  3.83197216]
```

```
In [90]: results = pd.DataFrame({
    'y_pred_test': y_pred_test,
    'y_test': y_test,
})
results
```

```
Out[90]:
```

	y_pred_test	y_test
15507	1.201223	1.270
96436	1.732650	1.804
139278	0.233131	0.302
151520	1.918767	1.982
38959	2.290398	2.376
...
104178	0.151963	0.216
13129	1.353629	1.396
77756	1.921646	2.008
88938	1.242944	1.314
95409	1.379472	1.424

55714 rows × 2 columns

```
In [91]: # C_values = [0.001, 0.01, 0.1, 1.0]
```


```
# mse_train_values = []
# mse_val_values = []
# # SVRModel = SVR(C=1.0, epsilon=0.1, kernel = 'rbf')
# for C in C_values:
#     model = SVR(C = C, epsilon=0.1, kernel = 'rbf')
#     mse_train, mse_val = eval_model(model, X_train, y_train, X_val, y_val)
#     mse_train_values.append(mse_train)
#     mse_val_values.append(mse_val)

# results = pd.DataFrame({
#     'C_values': C_values,
#     'mse_train': mse_train_values,
#     'mse_val': mse_val_values,
# })
```

```
In [92]: # for training set
y_pred_train = SVRModel.predict(X_train)
plt.scatter(y_train, y_pred_train, color='blue', label='Actual vs Predicted')
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)], linestyle='--', color='red', label='Perfect')


plt.title('Actual vs Predicted Values (Training Set)')
plt.xlabel('Actual Values')
```

```
plt.ylabel('Predicted Values')
plt.legend()
plt.show()
```



```
In [93]: # for validation set
y_pred_val = SVRModel.predict(X_val)
plt.scatter(y_val, y_pred_val, color='blue', label='Actual vs Predicted')
plt.plot([min(y_val), max(y_val)], [min(y_val), max(y_val)], linestyle='--', color='red', label='Perfect Prediction')

plt.title('Actual vs Predicted Values (Validation Set)')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.legend()
plt.show()
```



```
In [94]: # for test set
y_pred_test = SVRModel.predict(X_test)
plt.scatter(y_test, y_pred_test, color='blue', label='Actual vs Predicted')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], linestyle='--', color='red', label='Perfect Prediction')

plt.title('Actual vs Predicted Values (test Set)')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.legend()
plt.show()
```

